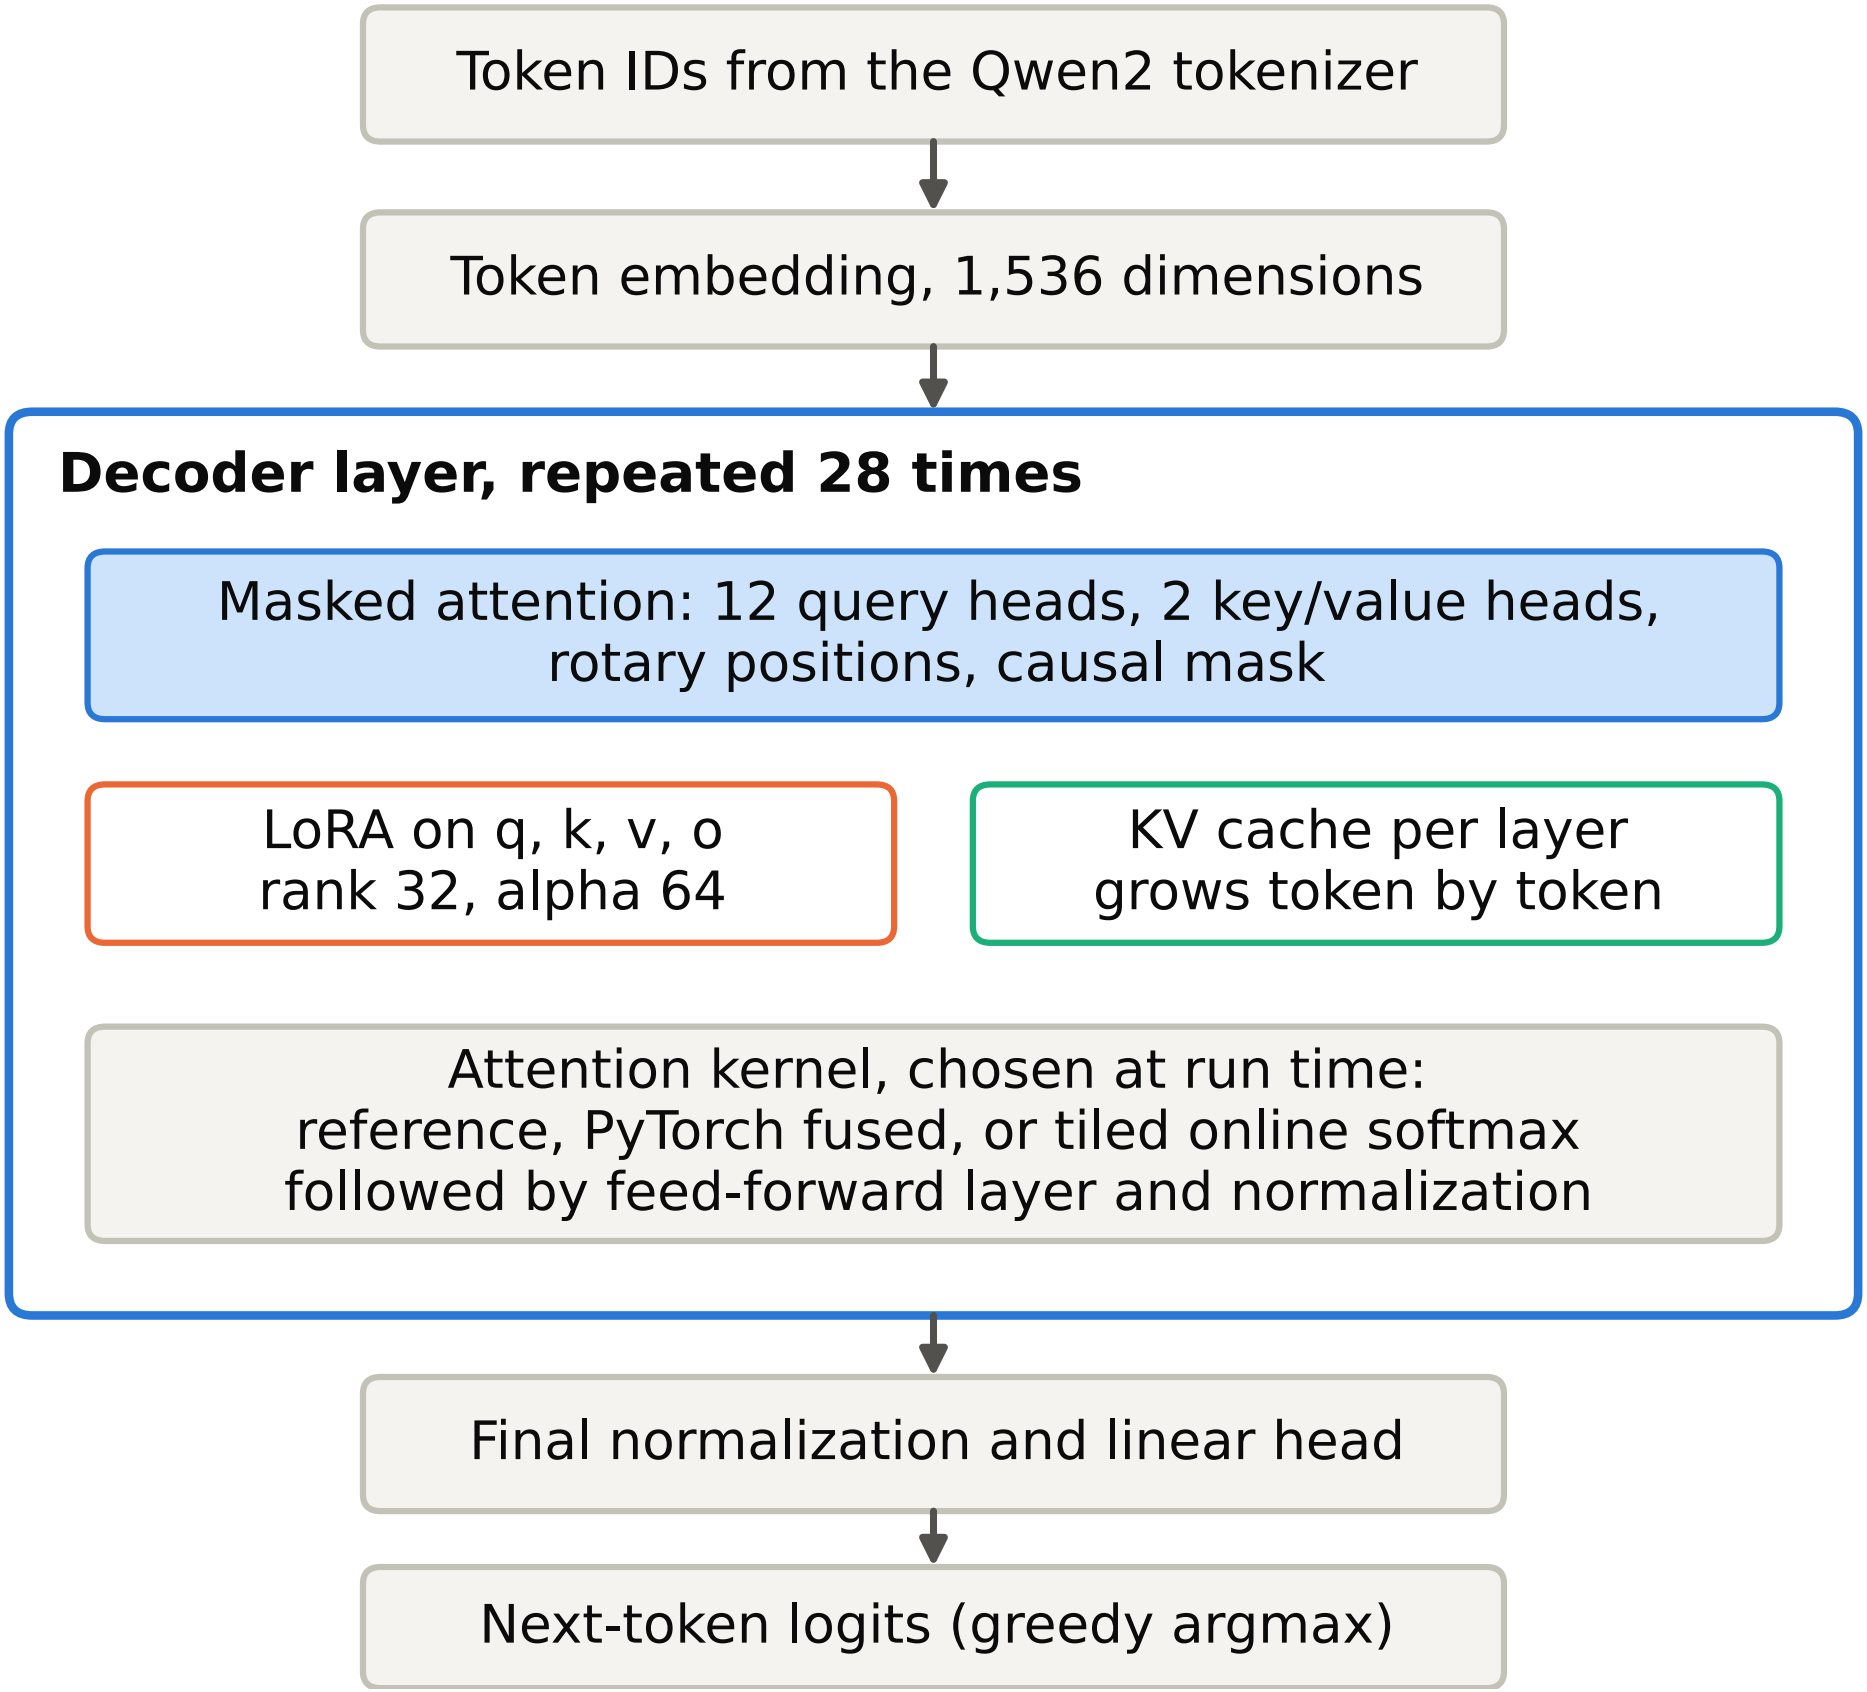
Token IDs from the Qwen2 tokenizer
Token embedding, 1,536 dimensions
Decoder layer, repeated 28 times
Masked attention: 12 query heads, 2 key/value heads, rotary positions, causal mask
LoRA on q, k, v, o rank 32, alpha 64
KV cache per layer grows token by token
Attention kernel, chosen at run time: reference, PyTorch fused, or tiled online softmax followed by feed-forward layer and normalization
Final normalization and linear head
Next-token logits (greedy argmax)