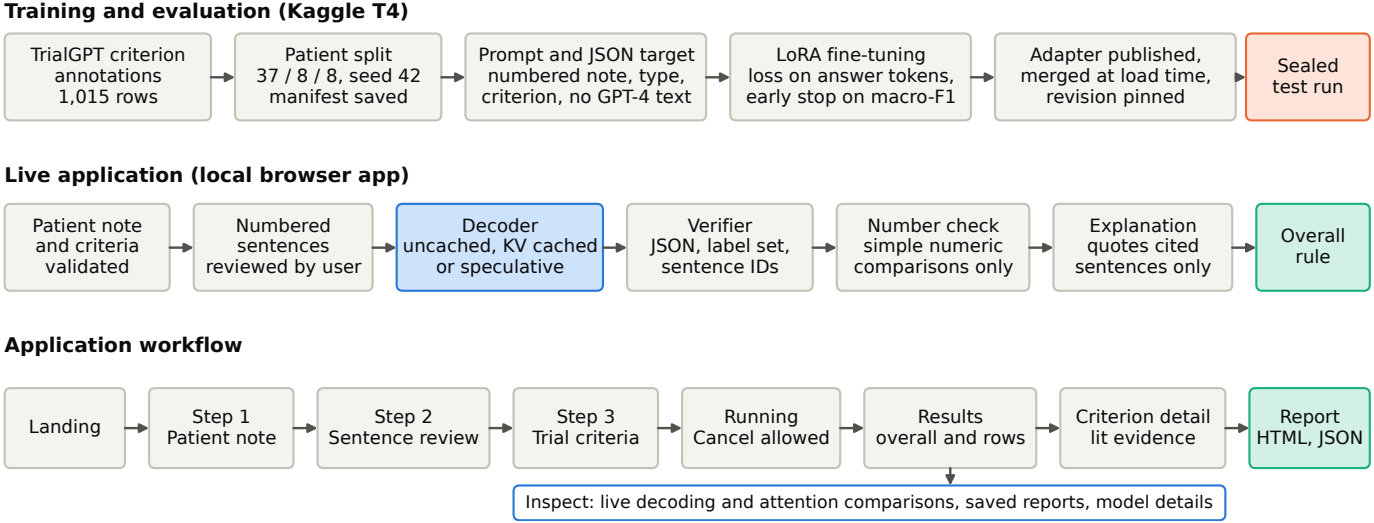
Training and evaluation (Kaggle T4)
TrialGPT criterion annotations 1,015 rows
Patient split 37 / 8 / 8, seed 42 manifest saved
Prompt and JSON target numbered note, type, criterion, no GPT-4 text
LoRA fine-tuning loss on answer tokens, early stop on macro-F1
Adapter published, merged at load time, revision pinned
Sealed test run
Live application (local browser app)
Patient note and criteria validated
Numbered sentences reviewed by user
Decoder uncached, KV cached or speculative
Verifier JSON, label set, sentence IDs
Number check simple numeric comparisons only
Explanation quotes cited sentences only
Overall rule
Application workflow
Landing
Step 1 Patient note
Step 2 Sentence review
Step 3 Trial criteria
Running Cancel allowed
Results overall and rows
Criterion detail lit evidence
Report HTML, JSON
Inspect: live decoding and attention comparisons, saved reports, model details